

House Price Prediction

This project builds a Random Forest model to predict housing prices using data from the Kaggle competition [House Prices: Advanced Regression Techniques](#). It demonstrates a full machine learning pipeline: data retrieval, preprocessing, model training, evaluation, and result visualization.

Project Files

File/Folder	Description
data_utils.py	Downloads and unzips data from Kaggle API
model.py	Trains a Random Forest model and evaluates it
images/	Stores output plots (prediction vs. actual, importance)
report.md	Final report for academic submission
README.md	This file

How to Run

1. **Place your Kaggle API key** in: ~/.kaggle/kaggle.json
2. **Download the dataset:**

```
python data_utils.py
```

3. **Train the model and generate plots:**

```
python model.py
```

4. **Check output in images/ folder:**

- pred_vs_actual.png

- `feature_importance.png`

Full Report

Read the full analysis and conclusions here: [report.md](#)

Requirements

Install required packages with:

```
pip install pandas numpy matplotlib scikit-learn kaggle
```